

Title: GARUDA AI: A Multi-Source, Cross-Verified, and Explainable AI Framework for Crime Pattern Analysis and Proactive Public Safety Intelligence

PROBLEM STATEMENT

Conventional crime prediction systems are largely reactive, typically rely on a single data modality, and where predictive models are used, they usually function as black boxes — producing a risk score with no justification and no independent verification of that score from more than one modelling angle. This lack of transparency and cross-verification limits trust and adoption among law enforcement agencies and policymakers. GARUDA AI addresses this by fusing structured crime records, spatial signals, and textual/social data through multiple, functionally distinct models, verifying each flagged prediction from statistical, temporal, and spatial perspectives, and explaining every output through several complementary explainability techniques before it reaches a decision-maker.

LITERATURE SURVEY

Eleven recent studies were reviewed across four themes central to GARUDA AI's design — explainable crime prediction, comparative/ensemble machine learning, spatio-temporal and clustering-based hotspot prediction, and NLP/social-media-based crime detection:

1	Ersöz, F., Ersöz, T., Marcelloni, F., & Ruffini, F. (2025)	Artificial Intelligence in Crime Prediction: A Survey With a Focus on Explainability — IEEE Access, 13	Reviews 142 AI crime-prediction studies; identifies trustworthiness/explainability as the field's central open problem.	Directly motivates GARUDA AI's explainability-first design.
2	Joshi, P., Sharma, R., & Verma, K. (2023)	Interpretable Machine Learning Models for Crime Prediction — ScienceDirect	Applies a single interpretable model type to crime data.	GARUDA AI extends this to multiple explainability techniques across multiple model families, not one model alone.
3	Mandalapu, V., Elluri, L., Vyas, P., & Roy, N. (2023)	Crime Prediction Using Machine Learning and Deep Learning: A Systematic Review and Future Directions — IEEE Access, 11	Systematic review of ML/DL crime-prediction approaches; flags lack of interpretability and cross-model verification.	Confirms the research gap GARUDA AI targets: verification across multiple models, not one.
4	Dubey, A., Venu, N., Bisen, D., Garg, P., & Srinivas, K. (2025)	A Clustering-based Method for Hotspot Recognition in Crime Forecasting — IEEE 8th ICICT	Uses clustering alone for hotspot recognition.	Used as the basis for GARUDA AI's spatial verification layer (DBSCAN/K-Means), extended with cross-checks from other models.
5	Vanitha, K., et al. (2025)	An Intelligent Crime Risk Prediction Framework using Behavioral Analysis and Advanced ML — ICSCSA 2025	Single-model behavioral-feature-based risk scoring.	Shows the common single-model risk-scoring pattern GARUDA AI moves beyond via multi-model consensus.
6	Mao, L., Du, W., Wen, S., Li, Q., Zhang, T., & Zhong, W. (2025)	Crime forecasting: A spatio-temporal analysis with deep learning models — SAGE journal	CNN-LSTM model for spatiotemporal crime-count forecasting.	Validates LSTM-based temporal modeling, used as GARUDA AI's temporal verification layer.
7	Anonymous authors (2025)	Hybrid Spatio-Temporal Crime Forecasting with Graph Neural Networks (DBSCAN + eLSTM + ST-GNN) — IIARD Int. J. of CS & Math Theory, 11(9)	Fuses DBSCAN clustering with LSTM and graph neural networks for forecasting.	Confirms clustering + temporal fusion improves accuracy, but adds no explainability or anomaly-based check — a gap GARUDA AI fills.
8	Butt, U.M., Letchmunan, S., Hassan, F.H., et al. (2020)	Spatio-Temporal Crime Hotspot Detection and Prediction: A Systematic Literature Review — IEEE Access, 8	Systematic review establishing clustering and time-series/deep-learning as standard hotspot techniques.	Establishes DBSCAN/K-Means as GARUDA AI's spatial-verification standard; notes the missing link to explainability.
9	Zapata Rozo et al. (2024)	An NLP-Based Framework to Spot Extremist Networks in Social Media — Complexity, Wiley, Article 3380488	NLP framework detecting hostile social manipulation signals for law enforcement.	Supports GARUDA AI's NLP Intelligence module for social-media threat-signal extraction.
10	Bokolo, B.G., Onyehanere, P., Ogegbene-Ise, E., Olufemi, I., &	Leveraging Machine Learning for Crime Intent Detection in Social Media Posts — AIGC 2023, Springer CCIS vol. 1946	Transformer-based ML detecting crime intent in social-media text.	Validates automated social-media crime-signal extraction as GARUDA AI's NLP module performs.

RESEARCH GAPS IDENTIFIED:

- Most existing systems use a single data modality — only structured records, only spatial data, or only text — rather than fusing all three into one pipeline.
- Where explainability is present, it is usually limited to one technique (SHAP or LIME alone) rather than multiple complementary methods used together.
- Predictions are rarely cross-verified from more than one modeling angle (statistical, temporal, spatial) before being reported as a final risk score.
- Evaluation in most studies stops at Accuracy or F1-Score; broader validation such as ROC-AUC, cross-validation, and multi-model consensus checking is uncommon.
- Very few frameworks progress from a research-stage model to a deployable, operational system with a live dashboard and alerting for actual police decision-making.

OBJECTIVES

1. To build a unified multi-source data pipeline integrating structured crime records, geospatial, weather/census, and textual (news/social media) data. **Technique:** ETL pipeline with automated cleaning, categorical encoding, geocoding, and normalization.
2. To develop a multi-model predictive core for crime risk estimation instead of relying on a single classifier. **Technique:** Ensemble of Logistic Regression, Decision Tree, Random Forest, XGBoost, LightGBM, and CatBoost.
3. To cross-verify every flagged high-risk prediction from independent statistical, temporal, and spatial perspectives before it is finalized. **Technique:** Isolation Forest (anomaly verification), LSTM (temporal verification), and DBSCAN/K-Means (spatial verification).
4. To ensure every prediction reaching a decision-maker is transparent, interpretable, and auditable. **Technique:** Multi-technique explainability layer — SHAP, LIME, Feature Importance, Partial Dependence, and Counterfactual/Anchor explanations. To deliver the resulting intelligence in an operational, decision-ready form rather than as a static research output. **Technique:** Interactive dashboard with hotspot maps and reports, plus a multi-channel (email/SMS/push) alert system.

NOVELTY / INNOVATION

1. To cross-verify high-risk predictions using statistical, temporal, and spatial perspectives before confirmation. **Technique:** Isolation Forest, LSTM, and DBSCAN/K-Means.
2. To confirm a location as high-risk only when independent verification agrees. **Technique:** Consensus-based gated decision layer combining outputs from all three verification models.
3. To provide explanations only for verified predictions rather than directly explaining unverified model outputs. **Technique:** Gated XAI layer triggered after consensus verification.
4. To provide multi-dimensional explanations for every confirmed high-risk prediction. **Technique:** SHAP, LIME, Feature Importance, Partial Dependence, and Counterfactual explanations.
5. To integrate verification and explainability into one unified pipeline, addressing the gap where existing approaches generally perform either prediction/verification or explainability separately. **Technique:** Prediction → Cross-Verification → Consensus Gate → Multi-Technique XAI → Decision Output

FEASIBILITY

1. To use publicly available benchmark datasets for development and evaluation.
Resources: UCI Communities and Crime, Chicago Crime, and NCRB Crime in India datasets.
2. To implement the proposed techniques using mature open-source libraries.
Tools: Scikit-learn, XGBoost, LightGBM, CatBoost, TensorFlow/PyTorch, SHAP, and LIME.
3. To minimize data-collection and preprocessing constraints by using well-documented datasets with existing research precedent.
Technique: Standardized data cleaning, transformation, normalization, and feature engineering.
4. To support computationally intensive model training within available resources.
Resources: Standard computing hardware or free-tier cloud platforms such as Google Colab/Kaggle.
5. To complete the complete pipeline within the proposed timeline using established algorithms, libraries, datasets, and documented implementation methods.
Pipeline: Data Processing → Prediction → Cross-Verification → Explainability → Dashboard.

PROJECT PLAN

Phase	Stage	Activities / Work to be Done
Phase 1	Conceptualization & Requirements Completed	• Define crime prediction problem • Conduct literature survey • Identify research gap and innovation • Define objectives and requirements • Collect and study benchmark datasets • Finalize methodology and system architecture
Phase 2	Data Engineering & Model Development	• Preprocess and integrate datasets • Feature engineering and data analysis • Train six core classifiers • Develop crime risk prediction models • Perform hotspot detection
Phase 3	Cross-Verification & Explainable Intelligence	• Implement Isolation Forest, LSTM, DBSCAN/K-Means • Build Cross-Verification Consensus Layer • Apply SHAP, LIME, Feature Importance, PDP, and Counterfactual explanations • Validate and compare model performance
Phase 4	Dashboard, Validation & Deployment	• Develop interactive dashboard • Integrate hotspot maps, risk scores, reports, and alerts • Perform complete performance evaluation • Deploy the system • Prepare documentation, research paper, and final demonstration

TEAM DETAILS

Team Members (Roll Numbers):

- 2420030672 – N. Sai Pragathi
- 2420030710 – S. Indu
- 2420030712 – A.V.N.S Anupama
- 2420030719 – T. Aashritha Reddy

Guide Name: Dr. K. Swapnika